


# Hosting Options

Infrastructure Comparison for Secure, Scalable, and Maintainable Deployment

June 21<sup>st</sup>, 2025 - Jan Schneider



# Hosting Needs of EXACT



## Compute

Execution of Docker containers for user submissions

Python-based workers (CPU-intensive, potential GPU support)

Low to moderate demand for backend (Django) and frontend (Next.js)




## Storage

At least 10 GB of storage for datasets and models (e.g., MRI, Tetris)

PostgreSQL database for users, leaderboards, and challenges

Temporary upload space for submitted explanation code



## Networking

Public HTTPS access (frontend and API)

Internal communication between containers (Docker Compose)

Secure API endpoints for challenge submission and administration



## Security & Maintainability

Container isolation for untrusted code execution (sandboxing)

User/admin authentication and access control

Optional: CI/CD, regular backups, system monitoring and logs

# Evaluation Criteria



## Cost

Is the solution affordable short- and long-term, are free credits available?



## Effort

How much time and expertise is required for setup, updates, and operation?



## Security

Does the option offer sufficient protection for data, code execution, and access?



## Scalability

Can the system grow with user demand and higher workloads?



# Overview



## AZURE

Academic-friendly, good AI tooling, Docker & PostgreSQL supported, GitHub Actions-native

Strong candidate due to free credits and ease of use

## AWS

Similar to AZURE, most flexible, huge ecosystem, but more complex (IAM, VPC, etc.)

Technically powerful, but higher setup overhead

## Self-Hosted

Full control, no external dependencies, runs on any server w/ Docker

Ideal for isolated environments, but high effort. Could be useful for early MVPs or isolated use cases



Overview of cloud-based architecture  
on Microsoft Azure

# Azure Infrastructure



## Frontend

## Backend

## Storage

## Support



Azure Static  
Web App



Azure  
Container  
App



Service Bus



Azure  
Container  
Job



Azure  
PostgreSQL



Azure Blob  
Storage



Key Vault



Insights

# Submission Processing Pipeline



## Support



Key Vault



Insights

Score fetch **8**

## Storage



Azure PostgreSQL

Save score & status **7**



Azure Blob Storage

## Backend



Azure Container App

**4** Trigger message to start evaluation



Service Bus

Start job with dataset, model, uploaded script **5**



Azure Container Job

Load script, dataset & model **6**

## Frontend

**1**

Submission request (form data + metadata)

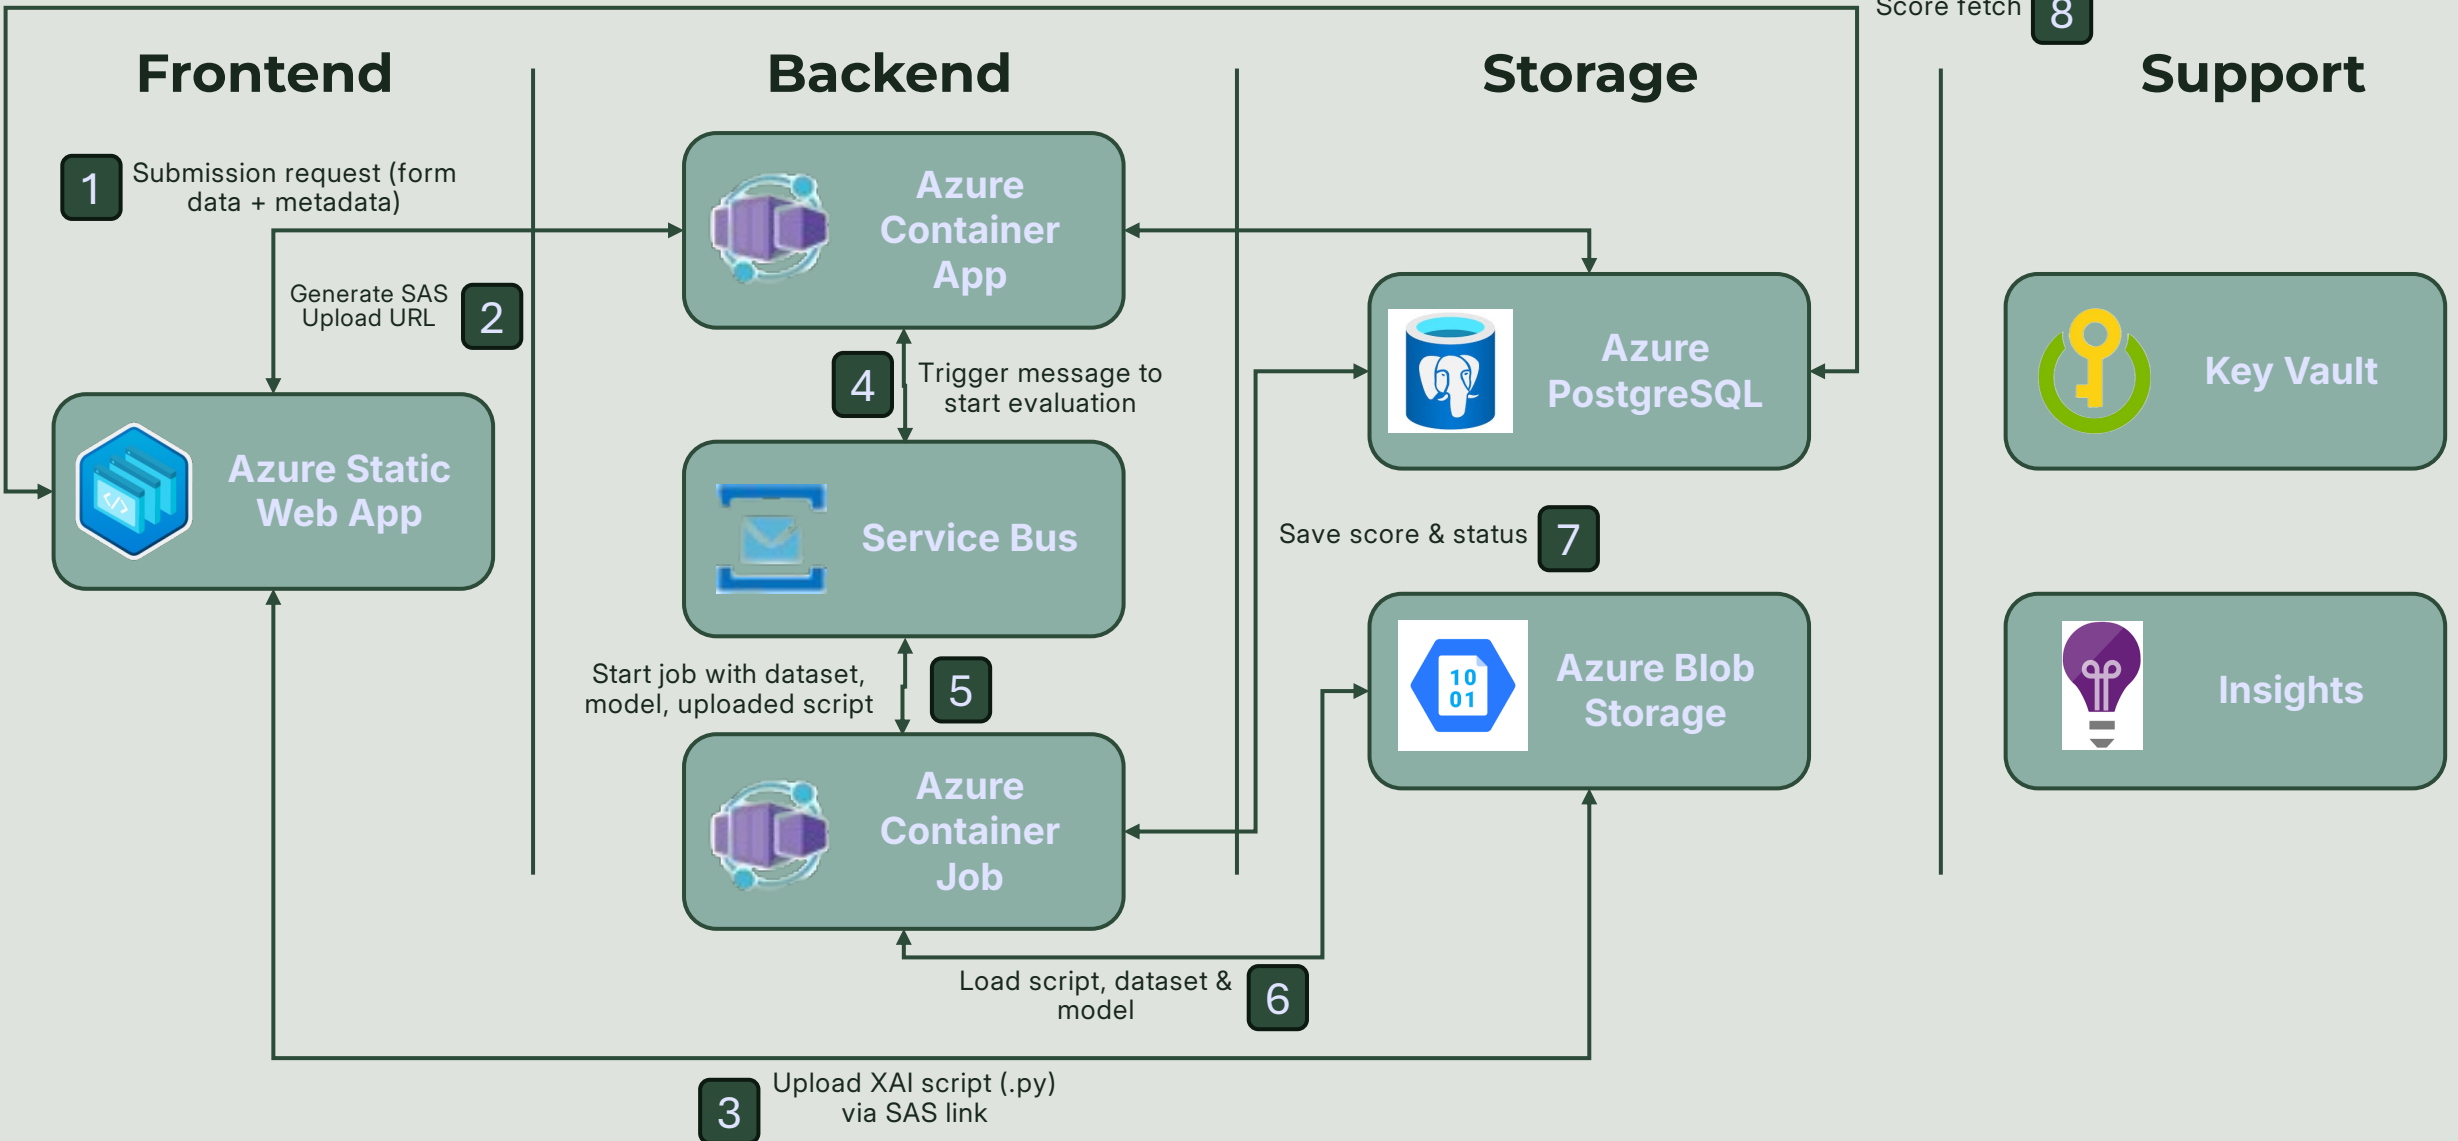
Generate SAS Upload URL **2**



Azure Static Web App

**3**

Upload XAI script (.py) via SAS link



# Conclusion Azure



## Advantages



### Scalability

Easily scale compute jobs and services using Azure Container Apps & Jobs



### Managed Infrastructure

No manual setup of databases, queues, or storage needed



### Security & Compliance

Azure offers built-in TLS, authentication, Key Vault integration



### Integrated Monitoring

Centralized insights via Azure Monitor and App Insights



### CI/CD Support

Seamless GitHub Actions integration for automatic deployment

## Disadvantages



### Cost Overhead

Higher baseline costs compared to self-hosted Docker setups



### Vendor Lock-In

Azure-specific features (e.g., SAS URLs, Service Bus) reduce portability



### Complexity

Managing permissions, services, and billing adds cognitive load



### Cold Start Latency

Container Jobs may take time to spin up on-demand



---

# Self-Hosted

Overview of Self-hosted infrastructure  
for transparent, flexible evaluation


# Self-Hosted Infrastructure



## Frontend



## Backend



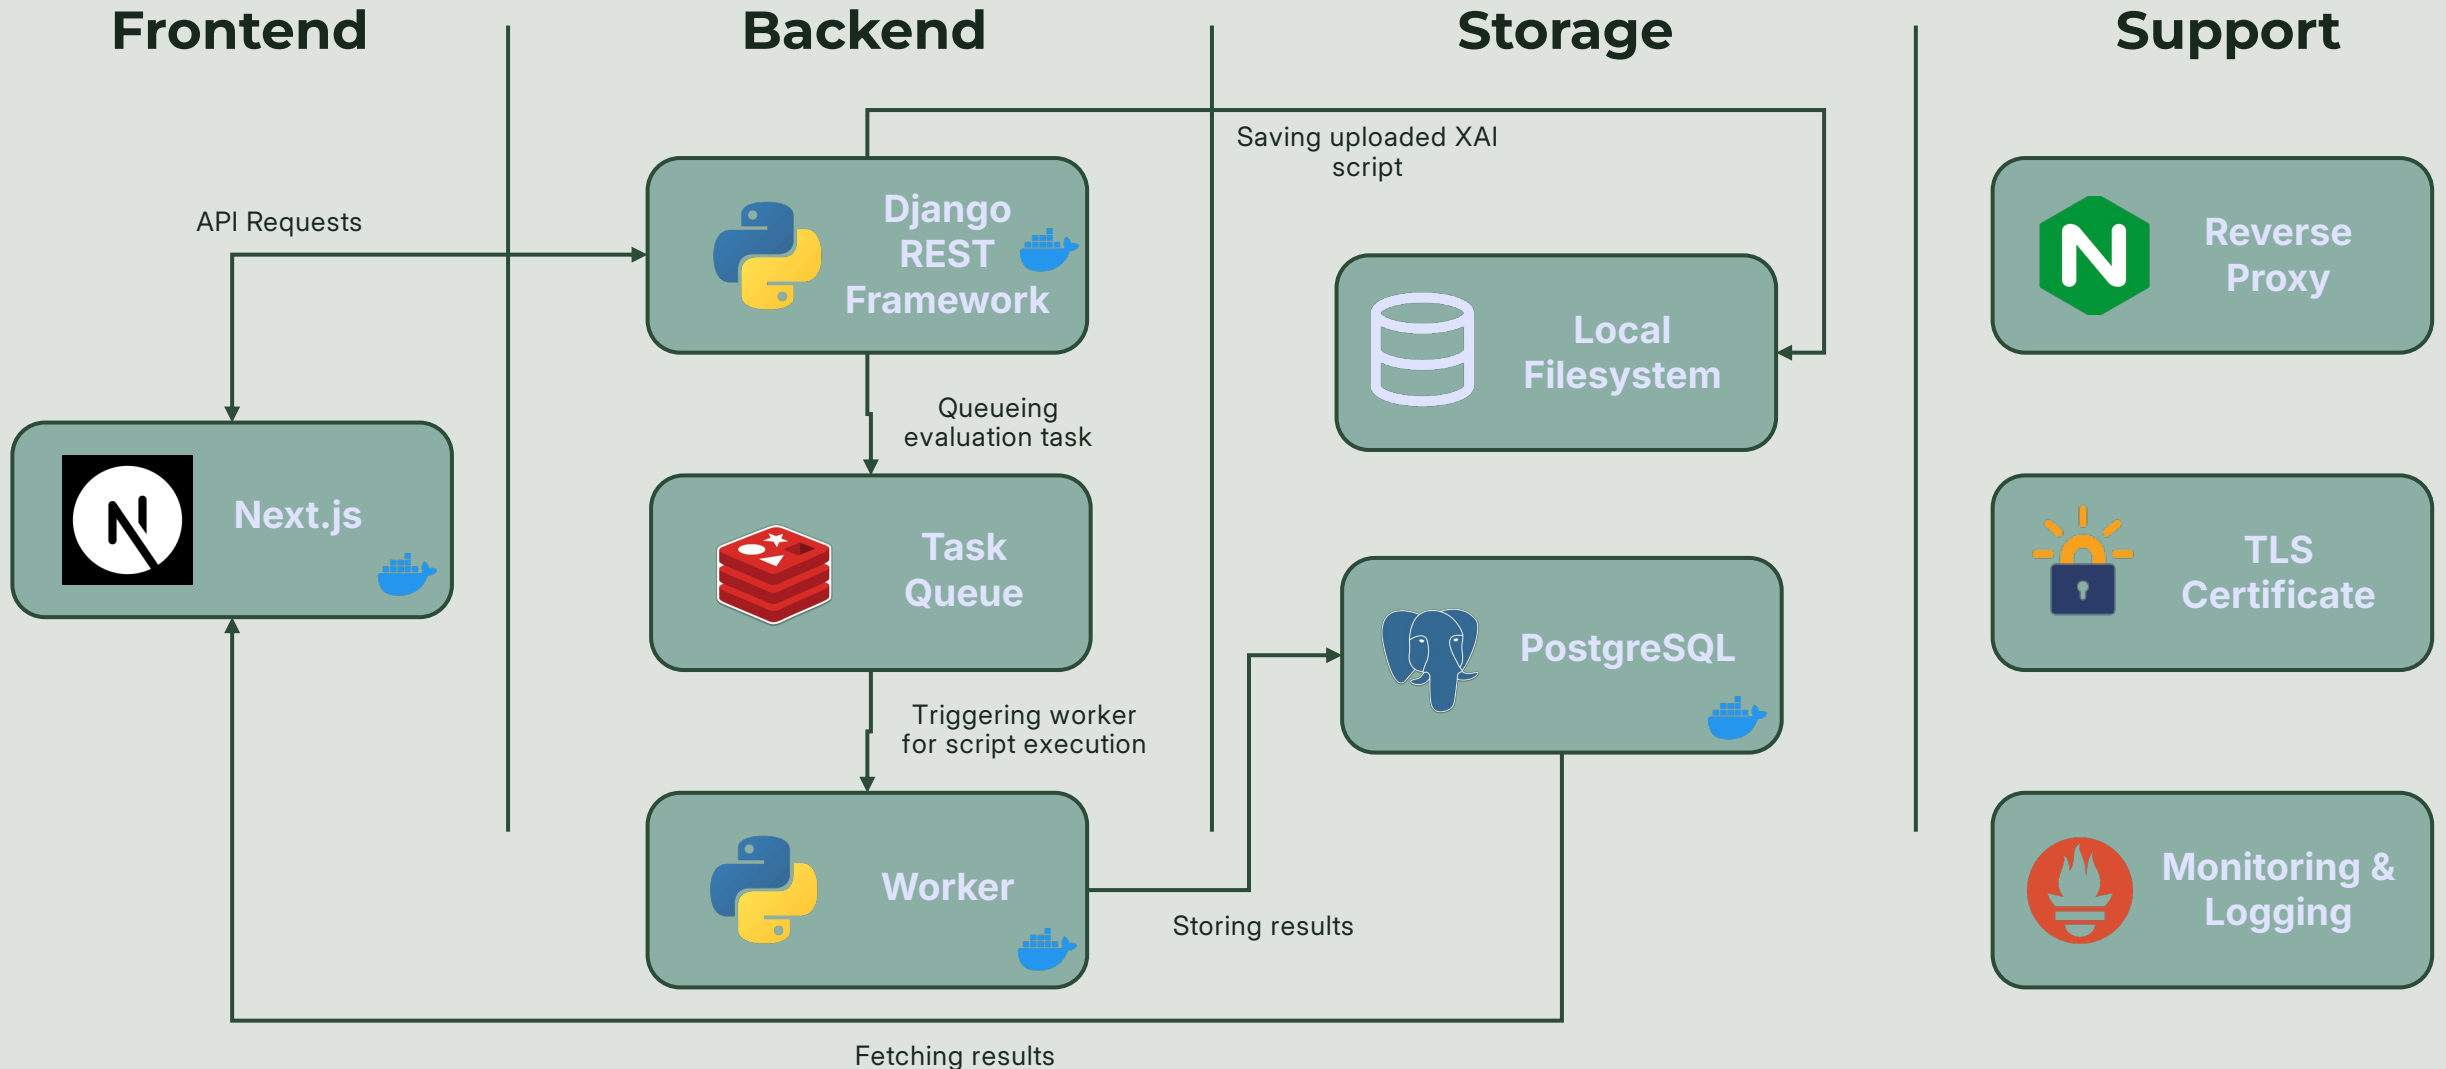
## Storage



## Support



# Self-Hosted Infrastructure



# Summary Self-Hosted



## Advantages



### Flexibility

Full control over architecture and services enables deep customization across all layers of the system



### Cost Efficiency

No platform or service fees; only pay for actual compute, storage, and traffic usage



### Tech Independence

Vendor-agnostic setup avoids lock-in and facilitates long-term maintainability



### Tooling Familiarity

Builds on established open-source tools such as Docker, PostgreSQL, and Redis

## Disadvantages



### Operational Overhead

Manual provisioning and maintenance of databases, job queues, and file storage



### Scalability Complexity

Horizontal scaling, high availability, and load balancing must be managed independently



### Security Responsibility

TLS, authentication, and secret management must be implemented and monitored in-house



### Monitoring & Observability

Requires integration of custom monitoring stacks (e.g., Prometheus, Grafana, Loki)

# Closing Statement



For professional or enterprise-grade deployments, I recommend using Azure (or alternatively AWS), as it offers built-in scalability, security, and monitoring – with minimal operational effort.



However, for academic settings, research prototypes, or environments where local server infrastructure is already available, a self-hosted setup can provide more flexibility and cost control – at the expense of increased maintenance responsibilities.